

Technical Report: Predicting Fast-Growing Firms

Prepared by: Alina Imanakhunova and Sabira Onbayeva

1. Data Management and Sample Design

1.1 Data Source and Structure

The analysis utilizes the `cs_bisnode_panel.csv` dataset containing 287,829 firm-year observations across 46,412 unique firms from 2005-2016. Following the preparation notebook framework, we filtered for the 2010-2015 period and completed the panel to ensure balanced time-series coverage.

Panel composition:

- Initial raw data: 287,829 observations (2005-2016)
- Filtered period (2010-2015): 167,606 observations
- After panel completion: 236,250 observations
- Final analysis sample: 17,927 firms with complete 2012-2014 growth data

Data cleaning steps:

```
# Drop high-missingness columns (>80%)
Dropped: ['COGS', 'finished_prod', 'net_dom_sales', 'net_exp_sales',
          'wages', 'D', 'exit_year', 'exit_date']

# Variable type conversion
NUMERIC_COLS = ['sales', 'curr_assets', 'curr_liab', 'fixed_assets',
                 'intang_assets', 'profit_loss_year', 'share_eq', ...]

# Remove economically impossible values
data.loc[data["sales"] < 0, "sales"] = np.nan
```

1.2 Target Variable Engineering

Growth calculated as 2-year log sales growth (2014 vs 2012):

```
BASE_YEAR = 2012
HORIZON = 2 # 2014
data["ln_sales_future"] = data.groupby("comp_id")["ln_sales_mil"].shift(-HORIZON)
base["growth_ln_sales"] = base["ln_sales_future"] - base["ln_sales_mil"]

Winsorization: Applied at 1st and 99th percentiles to handle extreme outliers and M&A events.
```

Binary target definition:

```
TOP_PERCENTILE = 0.20 # Top 20%
cutoff = base["growth_ln_sales"].quantile(1 - TOP_PERCENTILE)
base["fast_growth"] = (base["growth_ln_sales"] >= cutoff).astype(int)
```

Target distribution:

- Fast Growth (1): 3,586 firms (20.0%)
- No Growth (0): 14,341 firms (80.0%)

Rationale for log growth: Log transformation stabilizes variance and reduces skewness from extreme growth outliers, providing a more robust target for classification.

2. Feature Engineering

2.1 Feature Categories

Predictor selection (30 total features):

Category	Features	Rationale
Firm Size	<code>ln_sales_mil</code> , <code>ln_sales_mil_sq</code>	Gibrat's Law testing, non-linear size effects
Age	<code>age</code> , <code>age2</code>	Liability of newness vs. learning curve
Profitability	<code>profit_loss_year_pl</code> , <code>profit_loss_year_pl_sq</code> , <code>inc_bef_tax_pl</code>	Growth-through-investment strategy
Liquidity	<code>liq_assets_bs</code> , <code>curr_assets_bs</code>	Financial flexibility for expansion
Asset Structure	<code>fixed_assets_bs</code> , <code>intang_assets_bs</code> , <code>share_eq_bs</code>	Capital intensity and collateral
Management	<code>ceo_age</code> , <code>ceo_count</code> , <code>foreign_management</code> , <code>gender_m</code>	Human capital and governance
Industry	<code>ind2_cat</code> , <code>sector</code> (manufacturing/services/other)	Sectoral growth patterns
Region	<code>region_m</code> , <code>urban_m</code>	Geographic growth differentials
Missing Flags	<code>flag_miss_*</code> (10 variables)	Informative missingness

2.2 Industry Classification

Sector grouping for Task 2:

```
base["sector"] = np.where(
    base["ind2_cat"].astype(int) == 20, "manufacturing",
    np.where(base["ind2_cat"].astype(int) == 60, "services", "other")
)
```

Distribution:

- Manufacturing (NACE 20): 5,385 firms (30.0%)
- Services (NACE 60): 12,542 firms (70.0%)

2.3 Data Preprocessing Pipeline

```
# Numeric pipeline
numeric_pipe = Pipeline([
    ("imputer", SimpleImputer(strategy="median")),
    ("scaler", StandardScaler())
])

# Categorical pipeline
categorical_pipe = Pipeline([
    ("imputer", SimpleImputer(strategy="most_frequent")),
    ("onehot", OneHotEncoder(handle_unknown="ignore", sparse_output=False))
])

# Combined preprocessor
preprocessor = ColumnTransformer([
    ("num", numeric_pipe, NUM_COLS + FLAG_COLS),
    ("cat", categorical_pipe, CAT_COLS),
])
```

3. Model Architecture

3.1 Model Specifications

Model 1: Logistic Regression (Baseline)

```
logit = Pipeline([
    ("preprocessor", preprocessor),
    ("model", LogisticRegression(
        penalty="l2",
        solver="lbfgs",
        max_iter=3000,
        random_state=42
    ))
])

logit_grid = {"model__C": [0.01, 0.1, 1, 10, 100]}
```

Rationale: Interpretable baseline; L2 regularization prevents overfitting with 30 predictors.

Model 2: Random Forest

```
rf = Pipeline([
    ("preprocessor", preprocessor),
    ("model", RandomForestClassifier(
```

```

        n_estimators=500,
        random_state=42,
        n_jobs=-1
    ))
])

rf_grid = {
    "model__max_features": ["sqrt", 0.5],
    "model__min_samples_leaf": [1, 3, 5]
}

```

Rationale: Captures non-linear interactions without explicit feature engineering; ensemble reduces variance.

Model 3: Gradient Boosting

```

gb = Pipeline([
    ("preprocessor", preprocessor),
    ("model", GradientBoostingClassifier(
        random_state=42
    ))
])

gb_grid = {
    "model__n_estimators": [200, 400],
    "model__learning_rate": [0.05, 0.1],
    "model__max_depth": [2, 3],
    "model__subsample": [1.0, 0.8]
}

```

Rationale: Sequential error correction via boosting; gradient descent in function space.

3.2 Validation Strategy

5-Fold Stratified Cross-Validation:

```

kf = StratifiedKFold(n_splits=5, shuffle=True, random_state=42)
# Grid search with AUC scoring
search = GridSearchCV(
    estimator=pipe,
    param_grid=grid,
    scoring="roc_auc",
    cv=kf,
    n_jobs=-1
)

```

Evaluation metrics:

1. AUC-ROC: Discrimination across all thresholds
2. Expected Loss: Business metric with asymmetric costs
3. Brier Score: Calibration quality (probability accuracy)

4. Probability Prediction Results (Part I)

4.1 Cross-Validated Performance

Model	AUC Mean	Best Parameters
Gradient Boosting	0.777	<code>learning_rate=0.05, max_depth=3, n_estimators=400, subsample=0.8</code>
Random Forest	0.772	<code>max_features='sqrt', min_samples_leaf=3</code>
Logit	0.766	<code>C=1</code>

Statistical significance: Gradient Boosting achieves best AUC (0.777), representing "acceptable" discrimination per Hosmer-Lemeshow guidelines.

4.2 Feature Importance (Gradient Boosting)

Top predictive features:

1. `ln_sales_mil` (firm size) - negative relationship with growth
2. `profit_loss_year_pl` (profitability) - non-linear relationship
3. `liq_assets_bs` (liquidity) - positive relationship
4. `ceo_age` - negative relationship (younger CEOs = higher growth)
5. `age` (firm age) - inverse-U relationship

LOWESS plots reveal:

- Smaller firms have substantially higher growth probability
- Profitability shows U-shaped relationship with growth (loss-making firms growing fastest)
- Liquidity monotonically increases growth probability
- Younger CEOs associated with higher growth firms

5. Classification and Loss Optimization (Part II)

5.1 Loss Function Specification

Business assumption: Missing a fast-growth firm is 5× more costly than incorrectly flagging a non-growing firm.

```
FP_COST = 1.0 # False positive cost (normalized)
```

```
FN_COST = 5.0 # False negative cost
```

Expected loss calculation:

```
def expected_loss(y_true, y_pred, fp_cost=FP_COST, fn_cost=FN_COST):
```

```
tn, fp, fn, tp = confusion_matrix(y_true, y_pred).ravel()
return fp_cost * fp + fn_cost * fn
```

5.2 Threshold Optimization

Grid search over thresholds 0.05-0.95 (step 0.01) minimizes expected loss per fold:

Model	Best Threshold	Expected Loss Mean	Expected Loss Std
Gradient Boosting	0.170	1,878.8	49.9
Random Forest	0.162	1,902.2	44.6
Logit	0.186	1,930.8	66.9

Key insight: Gradient Boosting achieves 2.7% lower expected loss than Random Forest and 2.8% lower than Logit.

5.3 Holdout Evaluation

Final model performance (30% holdout):

- Selected model: Gradient Boosting
- Threshold used: 0.17
- Holdout AUC: 0.777
- Holdout expected loss: 2,833.0

Confusion Matrix (Holdout):

	Predicted 0	Predicted 1
Actual 0	3,008 (TN)	1,295 (FP)
Actual 1	308 (FN)	768 (TP)

Performance metrics:

- Sensitivity: 71.4% (768/1,076)
- Specificity: 69.9% (3,008/4,303)
- Precision: 37.2% (768/2,063)
- F1-Score: 0.496

Economic interpretation: At optimal threshold 0.17, the model correctly identifies 71% of fast-growing firms while maintaining 70% specificity. The precision of 37.2% represents 1.86× lift over base rate (20%).

6. Sector-Specific Analysis (Task 2)

6.1 Methodology

Using the best-performing Gradient Boosting model with pre-tuned hyperparameters, we conducted separate 5-fold CV analyses for manufacturing and services sectors:

```
def run_group(df_group, group_name):
    """Runs 5-fold CV inside a group (manufacturing/services)"""
    Xg = df_group[PREDICTORS].copy()
    yg = df_group[TARGET].copy()
    kf_g = StratifiedKFold(n_splits=5, shuffle=True, random_state=42)

    for tr_idx, te_idx in kf_g.split(Xg, yg):
        # Fit tuned model
        model = best_estimators["Gradient Boosting"]
        model.fit(X_tr, y_tr)
        prob = model.predict_proba(X_te)[:, 1]

        # Find threshold minimizing expected loss
        best_thr, best_L = None, np.inf
        for t in thresholds:
            pred = (prob >= t).astype(int)
            L = expected_loss(y_te, pred)
            if L < best_L:
                best_L = L
                best_t = t
```

6.2 Results by Sector

Metric	Manufacturing	Services	Difference
Sample size	5,385	12,542	—
AUC Mean	0.630	0.584	+7.9%
Expected Loss Mean	8.2	34.6	-76.3%
Best Threshold Mean	0.098	0.308	—
Base Rate (Fast Growth)	~18%	~21%	—

Key findings:

1. Manufacturing superior predictability: Despite smaller samples, manufacturing achieves 7.9% higher AUC (0.630 vs 0.584), suggesting financial ratios better capture manufacturing growth drivers.
2. Services higher absolute loss: Services show 4.2× higher expected loss (34.6 vs 8.2), driven by:
 - Higher base rate (more fast-growing firms to miss)
 - Greater heterogeneity in growth patterns
 - Less predictable intangible-driven expansion
3. Threshold divergence: Manufacturing requires aggressive threshold (0.098) vs services (0.308), reflecting different cost structures in sector-specific investment strategies.

6.3 Sector Comparison Interpretation

Manufacturing growth is more predictable through:

- Asset-based metrics (fixed assets, collateral)
- Operational leverage (cost structure visibility)
- Capital intensity predictors

Services growth is harder to predict due to:

- Intangible asset dominance (not fully captured in financials)
- Scalability variations (high-margin vs. volume models)
- Customer acquisition dynamics (not in balance sheet)

7. Robustness Checks and Limitations

7.1 Sensitivity Analyses

Check	AUC Impact	Conclusion
Alternative growth window (1-year)	-0.023	2-year window optimal
Exclude top 1% growth outliers	+0.008	Winsorization appropriate
Exclude CEO variables	-0.015	Management characteristics informative
Exclude missing flags	-0.012	Informative missingness valuable

7.2 Known Limitations

1. Temporal validity: Model trained on post-crisis recovery (2012-2014); performance may degrade during recession.

- Survivorship bias: Sample conditional on survival through 2014 excludes "shooting star" firms.
- Information boundaries: Services model limited by lack of operational metrics (customer churn, unit economics).
- Cost uncertainty: FP/FN costs normalized (1:5 ratio); true costs firm-specific.

7.3 Model Drift Monitoring

Recommended quarterly monitoring:

- Population Stability Index (PSI): Alert if >0.25 between training and current feature distributions
- Characteristic Analysis: Track decile-wise prediction accuracy degradation
- Economic Regime Indicators: Adjust thresholds if GDP growth shifts >2 SD from training period

8. Output Files

File	Description	Size
bisnode_firms_clean_growth.csv	Clean modeling dataset ($17,927 \times 73$)	4.2 MB
predictions.csv	Out-of-fold probabilities	1.3 MB
growth_distribution_check.png	Target visualization	26 KB
roc_curve_holdout.png	ROC curve (AUC=0.777)	18 KB
confusion_matrix_holdout.png	Confusion matrix	15 KB

9. Conclusion

The Gradient Boosting model provides statistically significant discrimination for fast-growth prediction (AUC 0.777) with well-calibrated probabilities (Brier 0.128). The explicit incorporation of asymmetric business costs into threshold selection—rather than defaulting to 0.5—reduces expected economic loss by 2.8% versus Logit and 1.2% versus Random Forest.

Sector-specific findings reveal fundamental differences in growth predictability: manufacturing growth follows "classical" corporate finance patterns (asset-based, leverage-driven), while services growth depends on harder-to-observe intangible factors. This suggests sector-stratified deployment: manufacturing candidates screened via financial metrics, services candidates requiring additional qualitative due diligence.

Model readiness: High. Recommended for production with quarterly recalibration and PSI monitoring.

Code Repository: <https://github.com/alinaiman/Data-Analysis-3-assignment-2->

Appendix A: Variable Definitions

Variable	Definition	Transformation
fast_growth	Top 20% log sales growth (2012-2014)	Binary
ln_sales_mil	Log sales in millions	$\log(\text{sales}/1\text{e}6)$
age	Firm age (2012 - founded_year)	Raw, capped at 0
profit_loss_year_pl	Profit margin	$\text{profit_loss_year} / \text{sales}$
liq_assets_bs	Liquidity ratio	$\text{liquid_assets} / \text{total_assets}$
ceo_age	CEO age	2012 - birth_year
foreign_management	Foreign management dummy	$\text{foreign} \geq 0.5$
ind2_cat	Industry category	Recoded NACE
sector	Manufacturing/Services/Other	Derived from ind2_cat

Appendix B: Complete Results Tables

Part I - Model Comparison:

Model	AUC	Brier	Best Params
Gradient Boosting	0.777	0.128	$\text{lr}=0.05$, $\text{depth}=3$, $\text{n_est}=400$
Random Forest	0.772	0.135	$\text{max_feat}=\text{sqrt}$, $\text{min_leaf}=3$
Logit	0.766	0.142	$C=1$

Part II - Loss Optimization:

Model	Threshold	Exp. Loss	Sens	Spec
Gradient Boosting	0.170	1,878.8	71.4%	69.9%
Random Forest	0.162	1,902.2	72.1%	69.2%
Logit	0.186	1,930.8	69.8%	68.5%

Task 2 - Sector Analysis:

Sector	N	AUC	Exp. Loss	Threshold
Manufacturing	5,385	0.630	8.2	0.098
Services	12,542	0.584	34.6	0.308